

Chapter 1

General Discussion

1.1 Summary of principal results

The work described within this thesis has centered around how prediction is used in sensory processes such as object recognition and prolonged learning. The work is heavily motivated by the LeabraTI (TI: Temporal Integration) framework (Chapter ??) which leverages the laminocolumnar structure of the neocortex (Mountcastle, 1997; Buxhoeveden & Casanova, 2002; Horton & Adams, 2005) to learn to predict temporally structured sensory inputs. Predictive learning in the LeabraTI framework is made possible by temporally interleaving predictions and sensory processing across the same populations of neurons so that powerful error-driven learning mechanisms (O'Reilly & Munakata, 2000; O'Reilly, Munakata, Frank, Hazy, & Contributors, 2012) can be used to compute a prediction error that can be learned against to minimize the difference between predictions and sensory events over time.

LeabraTI relies on a 10 Hz prediction-sensation period as its core “clock cycle”, suggested to correspond to the widely studied alpha rhythm observable across posterior cortex using scalp EEG (Palva & Palva, 2007; Hanslmayr, Gross, Klimesch, & Shapiro, 2011; VanRullen, Busch,

Drewes, & Dubois, 2011). Chapter ?? investigated the role of the alpha rhythm in prediction by using an entrainment paradigm (Schroeder, Lakatos, Kajikawa, Partan, & Puce, 2008; Calderone, Lakatos, Butler, & Castellanos, in press) in which stimuli were presented rhythmically at 10 Hz so that predictions and sensory information could be interleaved regularly at the optimal rate proposed by LeabraTI. The experiment made use of three-dimensional objects that required integration over multiple sequential views to extract their three-dimensional structure. Thus, relatively rapid predictive learning mechanisms that operate over subsequent 100 ms periods could be leveraged to optimally encode the the objects. The spatial coherence between views and temporal onset of each view were independently manipulated to determine their effect on stimulus encoding quality and the putative role of the alpha rhythm in predictive processing.

The results of the Chapter ?? experiment indicated that spatial coherence and predictable temporal onset of each stimulus in an entraining sequence enhanced discriminability of a subsequently presented probe stimulus. Oscillatory analyses indicated strong bilateral alpha power and phase coherence modulation as a function of stimulus predictability. Specifically, spatial predictability of entraining stimuli suppressed alpha power with a lower degree of phase alignment relative to unpredictable stimuli. Temporally predictable entraining stimuli had the opposite effect, with increased alpha power and phase alignment, indicating successful entrainment. Importantly, phase alignment due to temporal predictability remained elevated compared to temporally unpredictable stimuli during a 200 ms blank period between the entraining sequence and probe, indicating that the effects of temporal predictability could persist without exogenous entrainment. In addition to these bilateral main effects, right hemisphere sites exhibited synergistic effects of combined spatial and temporal probe predictability on EEG amplitude and 10 Hz phase coherence approximately 200 ms after probe onset.

Overall, the results of the Chapter ?? experiment support the basic claims put forward by the LeabraTI framework. The predictable 10 Hz presentation rate of the entraining sequence improved encoding of the target object, enhancing discriminability for the subsequent probe stimulus. This finding was accompanied by increased alpha phase alignment that remained elevated until the onset

of the probe, which was necessary for ensuring that the probe event was processed precisely when the brain was expecting sensory information and not when it was generating a prediction.

Given this basic support for the LeabraTI framework, the Chapter ?? experiment was designed to investigate the role of prolonged predictive learning of dynamic stimuli. Previous work has suggested that a temporal association rule similar to the one central to LeabraTI might be leveraged for constructing stable representations of spatially coherent visual inputs (Stringer, Perry, Rolls, & Proske, 2006; Wallis & Baddeley, 1997; Isik, Leibo, & Poggio, 2012) and indeed, a line of experiments by Di Carlo and colleagues demonstrated that predictable object transformation sequences build object invariance (Cox, Meier, Oertelt, & DiCarlo, 2005; Li & DiCarlo, 2008, 2010; Li & Dicarlo, 2012). Given these results, one might hypothesize that combined spatiotemporal predictability would be optimal for prolonged learning of three-dimensional objects.

The Chapter ?? experiment used a subset of the object stimuli from the previous chapter's experiment along with an explicit training period during which observers studied the objects while they were rotated were rotated through their views. The study period was followed by a series of test trials that required same-different judgements about static probe stimuli. Somewhat surprisingly, the results of the experiment were an almost complete reversal of the previous chapter's experiment. Discriminability was lowest when stimuli were learned in a combined spatiotemporally predictable context and highest when learned in a completely unpredictable context. Furthermore, there was some indication that the principal differences between predictability conditions during training were driven primarily by degenerate viewing angles caused by three-dimensional foreshortening in the objects used (Balas & Sinha, 2009; Farah, Rochlin, & Klein, 1994; Pizlo & Stevenson, 1999). In three out of four of the objects used in the experiment, accuracy was lower for degenerate views learned in a spatiotemporally predictable context compared to a completely unpredictable one.

Chapter ?? revisited the LeabraTI framework and described a neural network model that implemented the columnar substructure necessary for predictive learning. The model was trained to recognize the same three-dimensional objects used in the Chapter ?? and ?? experiments with

the goal of being able to reproduce the conflicting behavioral results of the experiments. LeabraTI predicts that spatially predictable sequences presented at a regular temporal interval should elicit a synergistic effect on behavioral measures due to the multiple prediction-sensation cycles that successfully integrate visuospatial information at optimal temporal intervals (see also Doherty, Rao, Mesulam, & Nobre, 2005; Rohenkohl, Gould, Pessoa, & Nobre, 2014). Such a synergistic effect was demonstrated in the Chapter ?? EEG results, but was simply additive for behavioral measures. Still, the model provided a reasonable account of these data. Furthermore, the model was able to produce the reversal effect observed in Chapter ?? by increasing the scale of a single projection of synaptic weights as a simple proxy for prolonged learning.

The model further indicated that the synaptic weight scaling that accompanies prolonged learning caused object invariance to “trickle down” to lower-level feature representations. This was problematic for the objects used, since some of them suffered from extreme foreshortening, causing severely degenerate views. This caused confusion between objects, potentially accounting for the overall reversal observed between the Chapter ?? and ?? experiments.

1.2 Open questions

1.2.1 Does the brain predict at 10 Hz or 5 Hz?

The results of the Chapter ?? experiment indicated that delta-theta band oscillations centered around 5 Hz also indexed predictability leading into the probe judgement. Specifically, 5 Hz power was suppressed due to spatial predictability during the 200 ms blank period between the entraining sequence and probe. 5 Hz power was also suppressed due to temporal predictability of the entraining sequence but then increased at the onset of the probe and remained elevated for over 250 ms after its appearance. 5 Hz phase angle alignment was elevated due to temporal predictability throughout the probe judgement and furthermore and exhibited synergistic enhancement in alignment when when spatial prediction was possible. Altogether, these results suggest delta-theta band oscillations, might also play a role in predictive processing in addition to the alpha oscillations

of primary interest.

There are two potential explanations for these delta-theta band effects. First, the 5 Hz effects could simply be a subharmonic of the fundamental effects of predictability observed at 10 Hz. Steady state visual evoked potentials (SSVEP) due to exogenous rhythmic stimulation are known to cause power increases at harmonic frequencies, in addition to the fundamental frequency of stimulation. It is unclear, however, whether increases in harmonic power are simply due to sharing the same zero crossing as the fundamental frequency waveform (and thus contributing power) or whether they actually might serve a functionally distinct role from the fundamental frequency (e.g., Herrmann, 2001; Kim, Grabowecy, Paller, & Suzuki, 2011).

The second potential explanation of the observed delta-theta band effects is that predictions also occur at 5 Hz. Two recent reviews have suggested that sensory prediction is at least partially subserved by delta-theta oscillations (Arnal & Giraud, 2012; Giraud & Poeppel, 2012). These theories were formulated to explain prediction during speech recognition, with a focus on multiple windows of integration that are necessary for robust comprehension of speech's hierarchy of time varying features (e.g., phonemes, syllables, words). Delta-theta band predictions support integration over periods of 150 ms or longer which can be used to predict the overall speech envelope (Aiken & Picton, 2008) or phrasal structure of utterances. These theories are supported by recent empirical evidence of the importance of delta-theta band oscillations in the speech (Arnal, Wyart, & Giraud, 2011) and general audition domains (Stefanics et al., 2010). Others have noted, however, that auditory processing does not exhibit the ubiquitous perceptual discretization that visual processing does (VanRullen, Zoefel, & Ilhan, 2014). This does not necessarily mean that audition is not a predictive process. It could suggest, however, that auditory prediction happens at a more abstract level after feature extraction, consistent with the slower overall prediction rate, opposed to the sensory level for visual prediction and as suggested by LeabraTI.

In the Chapter ?? experiment, 5 Hz power and phase coherence were sensitive to the spatial predictability of the entraining sequence, and so it seems unlikely that they were simply a harmonic side effects of 10 Hz rhythmic stimulation. Thus, the most reasonable explanation of the 5 Hz

predictability effects observed in the Chapter ?? experiment results is that they reflected a slower, more abstract visual prediction process such as anticipation of whether the probe depict the same object as the entraining sequence or a different object.

1.2.2 Are “paper clip” objects somehow special?

1.2.3 Is synaptic weight scaling a good proxy for prolonged learning?

1.3 Conclusion

References

- Aiken, S., & Picton, T. (2008). Human cortical responses to the speech envelope. Ear and Hearing, 29(2), 139–157.
- Arnal, L. H., & Giraud, A.-L. (2012). Cortical oscillations and sensory predictions. Trends in Cognitive Sciences, 16(7), 390–398.
- Arnal, L. H., Wyart, V., & Giraud, A.-L. (2011). Transitions in neural oscillations reflect prediction errors generated in audiovisual speech. Nature Neuroscience, 14(6), 797–801.
- Balas, B. J., & Sinha, P. (2009). The role of sequence order in determining view canonicity for novel wire-frame objects. Attention, Perception & Psychophysics, 71(4), 712–723.
- Buxhoeveden, D. P., & Casanova, M. F. (2002). The minicolumn hypothesis in neuroscience. Brain, 125(Pt 5), 935–951.
- Calderone, D. J., Lakatos, P., Butler, P. D., & Castellanos, F. X. (in press). Entrainment of neural oscillations as a modifiable substrate of attention. Trends in Cognitive Sciences.
- Cox, D. D., Meier, P., Oertelt, N., & DiCarlo, J. J. (2005). 'Breaking' position-invariant object recognition. Nature Neuroscience, 8(9), 1145–1147.
- Doherty, J. R., Rao, A., Mesulam, M. M., & Nobre, A. C. (2005). Synergistic effect of combined temporal and spatial expectations on visual attention. The Journal of Neuroscience, 25(36), 8259–8266.
- Farah, M., Rochlin, R., & Klein, K. (1994). Orientation invariance and geometric primitives in shape recognition. Cognitive Science, 18(2), 325–344.
- Giraud, A.-L., & Poeppel, D. (2012). Cortical oscillations and speech processing: Emerging computational principles and operations. Nature Neuroscience, 15(4), 511–517.
- Hanslmayr, S., Gross, J., Klimesch, W., & Shapiro, K. L. (2011). The role of oscillations in temporal attention. Brain research reviews, 67(1-2), 331–343.
- Herrmann, C. S. (2001). Human EEG responses to 1-100 Hz flicker: Resonance phenomena in visual cortex and their potential correlation to cognitive phenomena. Experimental Brain Research, 137(3-4), 346–353.

- Horton, J. C., & Adams, D. L. (2005). The cortical column: A structure without a function. Philosophical Transactions of the Royal Society B, 360(1456), 837–862.
- Isik, L., Leibo, J. Z., & Poggio, T. (2012). Learning and disrupting invariance in visual recognition with a temporal association rule. Frontiers in Computational Neuroscience, 6(37).
- Kim, Y.-J., Grabowecky, M., Paller, K. A., & Suzuki, S. (2011). Differential roles of frequency-following and frequency-doubling visual responses revealed by evoked neural harmonics. Journal of Cognitive Neuroscience, 23(8), 1875–1886.
- Li, N., & DiCarlo, J. J. (2008). Unsupervised natural experience rapidly alters invariant object representation in visual cortex. Science, 321(5895), 1502–1507.
- Li, N., & DiCarlo, J. J. (2010). Unsupervised natural visual experience rapidly reshapes size-invariant object representation in inferior temporal cortex. Neuron, 67(6), 1062–1075.
- Li, N., & Dicarlo, J. J. (2012). Neuronal learning of invariant object representation in the ventral visual stream is not dependent on reward. The Journal of Neuroscience, 32(19), 6611–20.
- Mountcastle, V. B. (1997). The columnar organization of the neocortex. Brain, 120(Pt 4), 701–722.
- O'Reilly, R. C., & Munakata, Y. (2000). Computational Explorations in Cognitive Neuroscience: Understanding the Mind by Simulating the Brain. Cambridge, MA: The MIT Press.
- O'Reilly, R. C., Munakata, Y., Frank, M. J., Hazy, T. E., & Contributors (2012). Computational Cognitive Neuroscience. Wiki Book, 1st Edition, URL: <http://ccnbook.colorado.edu>.
- Palva, S., & Palva, J. M. (2007). New vistas for alpha-frequency band oscillations. Trends in Neurosciences, 30(4), 150–158.
- Pizlo, Z., & Stevenson, A. K. (1999). Shape constancy from novel views. Perception & Psychophysics, 61(7), 1299–1307.
- Rohenkohl, G., Gould, I. C., Pessoa, J., & Nobre, A. C. (2014). Combining spatial and temporal expectations to improve visual perception. Journal of Vision, 14(4), 1–13.
- Schroeder, C. E., Lakatos, P., Kajikawa, Y., Partan, S., & Puce, A. (2008). Neuronal oscillations and visual amplification of speech. Trends in Cognitive Sciences, 12(3), 106–113.
- Stefanics, G., Hangya, B., Herndi, I., Winkler, I., Lakatos, P., & Ulbert, I. (2010). Phase entrainment of human delta oscillations can mediate the effects of expectation on reaction speed. The Journal of Neuroscience, 30(41), 13578–13585.
- Stringer, S. M., Perry, G., Rolls, E. T., & Proske, J. H. (2006). Learning invariant object recognition in the visual system with continuous transformations. Biological Cybernetics, 94(2), 128–142.
- VanRullen, R., Busch, N. A., Drewes, J., & Dubois, J. (2011). Ongoing EEG phase as a trial-by-trial predictor of perceptual and attentional variability. Frontiers in Psychology, 2(60).
- VanRullen, R., Zoefel, B., & Ilhan, B. (2014). On the cyclic nature of perception in vision versus audition. Philosophical Transactions of the Royal Society B, 369(1641).
- Wallis, G., & Baddeley, R. (1997). Optimal, unsupervised learning in invariant object recognition. Neural Computation, 9(4), 883–894.