

Chapter 1

The LeabraTI framework: Spatiotemporal prediction with thalamocortical rhythms

1.1 Introduction

This chapter describes the LeabraTI (Temporal Integration) framework, which is a mechanistic description and general model of how prediction and temporal integration works in the brain. It is closely related to the Simple Recurrent Network (SRN) (Elman, 1990; Servan-Schreiber, Cleeremans, & McClelland, 1991) a neural network architecture that explicitly represents temporally lagged information in discrete “context” units whose activity gets integrated with more current information to predict what happens in the next time step (Figure 1.1A). This method of copying a contextual representation from an intermediate representation at discrete intervals was originally shown to be a robust way to leverage error-driven learning to represent latent temporal structure in auditory streams and artificial grammars. More generally, the SRN’s explicit representation of temporally lagged context can capture the latent structure of any stimulus that varies systematically over time, making it a good basis for a generic prediction and temporal integration mechanism.

LeabraTI differs in several key ways from the classical SRN architecture, primarily in the way context is represented and used in predictive processing. These differences are due to biologi-

cal constraints imposed by the microcircuitry of the neocortex, and thus form a number of testable predictions that can be used to evaluate the validity of the LeabraTI framework. The central prediction of LeabraTI is that temporally lagged context is represented by deep (Layer 6) neurons, which is possible in part to the bifurcation of intra-areal and inter-areal processing streams. As neural processing is a continuous operation, LeabraTI requires a regular interval over which to integrate deep context and make predictions, which is approximately every 100 ms. Predictions are made by driving superficial (Layers 2 and 3) neurons with the state of deep neurons through the intra-areal pathway, which is interleaved with standard peripheral sensory inputs over a total period that is also 100 ms. The strong 100 ms dependency in LeabraTI corresponds to the brain's alpha rhythm, which has been studied extensively using scalp EEG (Mathewson, Gratton, Fabiani, Beck, & Ro, 2009; Busch, Dubois, & VanRullen, 2009; Gould, Rushworth, & Nobre, 2011; Rohenkohl & Nobre, 2011; Mathewson et al., 2012; Belyusar et al., 2013; Palva & Palva, 2007; Hanslmayr, Gross, Klimesch, & Shapiro, 2011; VanRullen & Dubois, 2011).

TODO: Foreshadow structure of chapter...

1.2 LeabraTI biological details

1.2.1 Laminar structure and microcircuitry of the neocortex

A salient feature of the brain, and potential clue in realizing how an SRN-like computation might be carried out in biological neural circuits, is the laminar structure prevalent across the neocortex (Figure 1.1B). Incoming information from the sensory periphery is transmitted through the thalamus and targets Layer 4 neurons in the primary sensory cortices (e.g., V1). From there, Layer 4 neurons propagate spikes to superficial neurons (Layers 2 and 3) which in turn target Layer 4 neurons of higher-level cortices, forming the prominent corticocortical feedforward pathways that subserve visual and auditory recognition (Felleman & Van Essen, 1991). Corticocortical feedback originates in superficial layers or Layer 6 of the higher-level cortex and generally terminates on superficial neurons of the lower-level cortex (Rockland & Pandya, 1979). In addition to these

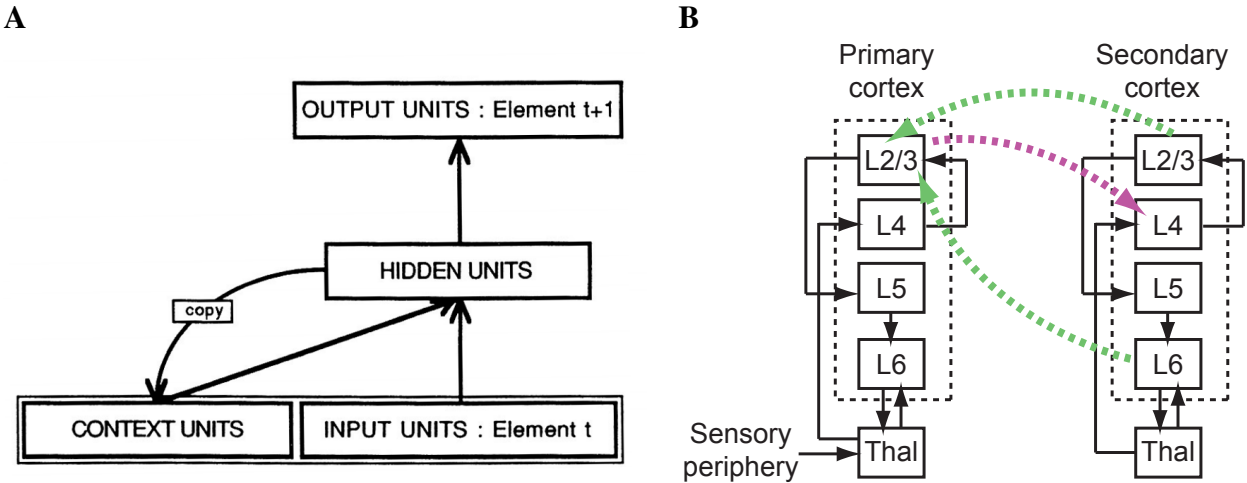


Figure 1.1: The Simple Recurrent Network (SRN) and microcircuitry of the neocortex

A: The SRN represents temporal information explicitly using discrete context units that are updated once per time step. Context is integrated with more current inputs to predict information at the subsequent time step. Reproduced from Servan-Schreiber et al. (1991). **B:** The neocortex is laminated with canonical circuitry between neurons across layers and between areas. Principal intra-areal connections are shown in black with inter-areal feedforward connections in purple and feedback connections in green.

inter-areal pathways, there exists a canonical microcircuit of the form Layer 4 \rightarrow Layer 2/3 \rightarrow Layer 5 \rightarrow Layer 6 that routes spike propagation through the local neuronal structure (Douglas & Martin, 2004; Thomson & Lamy, 2007; da Costa & Martin, 2010). This microcircuit forms the core computational unit of LeabraTI, as will be described in this and the following sections.

The importance of the local microcircuit was first suggested by Vernon Mountcastle in his proposal regarding the gross columnar organization of the neocortex (see Mountcastle, 1997, for a comprehensive review). Mountcastle's proposal states that microcolumns composed of around 80-100 neurons extending vertically through all six lamina with canonical circuitry form the core repeating structure of the neocortex. Neurons within a single microcolumnar circuit possess nearly identical receptive field tunings across lamina while neurons in neighboring microcolumns (radial separation greater than 600 μm) possess very different receptive field tunings but contribute to the higher-order macrocolumn (i.e., hypercolumn) structure (Hubel & Wiesel, 1977; Jones, 2000). Microcolumns have been identified in a variety of neural systems with this electrophysiological

mapping and are also prominently visible under Nissl staining. Despite this evidence for their structural existence, any function of the microcolumn aside from an organizing principle remains debated (Buxhoeveden & Casanova, 2002; Horton & Adams, 2005).

LeabraTI provides a computational role for the microcolumn, by mapping an SRN-like computation onto their Layer 4 \rightarrow Layer 2/3 \rightarrow Layer 5 \rightarrow Layer 6 circuit (Figure 1.1). In this mapping, superficial neurons continuously integrate feedforward and feedback inter-areal synapses to process current information. Layer 2/3 \rightarrow Layer 5 \rightarrow Layer 6 provides an intra-areal pathway for explicitly representing temporal context in Layer 6 neurons, which are relatively isolated from nonlocal inputs. There is also appropriate circuitry for recirculating this context through the local microcolumn via Layer 4 to drive the learning of temporal associations. This basic idea provides a concise explanation for the strong degree of isotuning throughout a single microcolumn, as Layer 6 neurons need to represent the same overall information as superficial neurons except at a delayed interval.

1.2.2 Layer 5 rhythmic bursting and contextual gating

The laminocolumnar organization of the neocortex provides the dual pathways necessary for continuous information processing and the SRN's explicit temporal context representation. The Layer 4 \rightarrow Layer 2/3 \rightarrow Layer 5 \rightarrow Layer 6 microcircuit only contains four synapses plus the transthalamic re-entrant synapses. Intracolumnar monosynaptic latencies for regular spiking neurons are on the order of 5 ms or faster (Armstrong-James, Fox, & Das-Gupta, 1992; Lumer, Edelman, & Tononi, 1997) and thus this relatively small amount of tissue, if driven with constant input, would circulate spikes at a rate too fast to perform enough temporal integration to make useful predictions. Several studies have noted that a subset of Layer 5 neurons exhibit intrinsic bursting at ~ 10 Hz when over threshold (Connors, Gutnick, & Prince, 1982; Silva, Amitai, & Connors, 1991; Franceschetti et al., 1995). This rhythmic bursting might implement a gating mechanism for updating Layer 6 context information at a regular 100 ms interval.

More specifically, Layer 5 neurons can be roughly divided into 5a and 5b subtypes (Thom-

son & Lamy, 2007). Layer 5a neurons have relatively small cell bodies and exhibit regular spiking depolarization responses. They collect input from other Layer 5a neurons both within and across columns (Schubert, Kotter, & Staiger, 2007) and pass it to 5b neurons and thus, likely play a simple information integration role. Layer 5b neurons, in contrast, have larger cell bodies and exhibit the aforementioned 10 Hz intrinsic bursting response profile. In the context of LeabraTI, the interpretation of this data is that the 5a neurons serve to integrate information from multiple Layer 2/3 neurons, with the 5b neurons gating context to Layer 6 neurons with each 10 Hz burst.

Layer 6 corticothalamic neurons receive strong inputs from Layer 5b neurons and send axons toward the thalamus completing the microcircuit within the local column and allowing the temporally lagged Layer 6 responses to integrate with more current Layer 4 inputs. Information is relayed from the thalamus back up to layer 4 in a focal one-to-one manner that maintains microcolumnar separation (Sherman & Guillery, 2006; Thomson, 2010), which could allow temporal associations to be formed by local Hebbian learning mechanisms that track high probability co-occurrences across past and present events (Foldiak, 1991).

1.2.3 Thalamic gating and sensory prediction

Both the SRN computation and the Leabra algorithm (O'Reilly & Munakata, 2000; O'Reilly, Munakata, Frank, Hazy, & Contributors, 2012) that are used to implement the LeabraTI framework are predicated on using powerful error-driven learning mechanisms (in addition to more standard Hebbian learning mechanisms) to represent the mapping between sensory inputs and outputs. In the context of temporal integration, error-driven learning would allow computation of error signals based on the difference between what is predicted to happen at a given moment (given the previous moments context as an input) and what actually happens. However, this computation requires that both the prediction and the actual sensation are represented by the same neural tissue so that an error signal can be computed, which is not possible if the sensory periphery is continuously transmitting incoming inputs.

To resolve this issue, the LeabraTI framework posits that predictions about sensory events

and the sensory events themselves are temporally interleaved through the same population of neurons in an alternating manner. This requires a mechanism to periodically downmodulate or even block the transmission of inputs from the sensory periphery. A subset of cells in the thalamus exhibit ~ 10 Hz intrinsic bursting properties similar to that of Layer 5 neurons (Lopes da Silva, 1991; Hughes et al., 2004; Lorincz, Crunelli, & Hughes, 2008; Lorincz, Kekesi, Juhasz, Crunelli, & Hughes, 2009), and thus perhaps perform a similar gating computation of sensory inputs into cortical circuits. In the context of LeabraTI, these bursting neurons might shift the balance of inputs to Layer 4 and superficial neurons between endogenous inputs local to the microcolumn representing predictions and quick bursts of actual sensory information.

More specifically, during the non-bursting intervals of thalamic intrinsically bursting neurons' response, environmental inputs are downmodulated due to the relative quiescence and Layer 6 neurons provide the dominant driving potential to the microcolumn. Layer 6 corticothalamic neurons exhibit a strong regular spiking depolarization response with facilitating short-term dynamics (Thomson, 2010) unlike all other pyramidal neurons, which exhibit depressing dynamics. This might suggest a specialized function for Layer 6 neurons, which in the context of LeabraTI is to drive a sustained prediction about upcoming sensory information. These layer corticothalamic neurons also sustain their drive through reciprocal projections with the thalamic relay cells that they project to. In addition to the Layer 6 \rightarrow Layer 4 transthalamic pathway, Layer 6 neurons also project directly to Layer 4. While these projections are relatively weak (Hirsch & Martinez, 2006), they do activate a metabotropic glutamate receptor (mGluR) that produces sustained depolarization similar to Layer 6 corticothalamic neurons (Lee & Sherman, 2009) - these direct ascending synapses are another possible route for sustained context information to drive Layer 4 neurons.

The burst response of thalamic intrinsic bursting neurons synchronized with the burst response of Layer 5b neurons destabilizes the Layer 6 sustained prediction through biased competition. This provides a snapshot of the current state of the sensory periphery as well as an opportunity to integrate the state of current sensory event into a prediction about what will happen during the next moment. At this time, standard error-driven learning mechanisms that compute

short timescale firing rate differences (O'Reilly & Munakata, 2000; O'Reilly et al., 2012) compute an error signal between the previous moment's prediction and the sensory outcome to minimize overall prediction error.

1.3 Summary of LeabraTI computation

The overall computation of LeabraTI is shown in Figure 1.2 and summarized here. Standard Leabra processing operates across two phases: a *minus* phase that represents the system's expectation for a given input and a *plus* phase, representing observation of the outcome. Generally, feedforward sensory inputs provide the dominant drive during the minus phase while the plus phase is driven by a combination of feedforward sensory inputs and the target outcome, conveyed via feedback projections. Leabra networks typically only model at the level of superficial (Layers 2 and 3) neurons, since they provide the principal feedforward and feedback projections necessary for error-driven learning.

In LeabraTI, deep Layer 6 neurons provide the dominant input to superficial Layer 2/3 neurons during the minus phase, relayed via Layer 4 transthalamic and direct ascending projections (not pictured in Figure 1.2). This represents a sensory prediction about what is about to happen during the subsequent plus phase, using temporally lagged context from the previous plus phase. To be computationally useful, the prediction needs to last at least 50 ms in order to allow contributions from both the feedforward drive and slightly longer latency feedback from higher-order cortical areas. The short-term facilitating dynamics unique to Layer 6 corticothalamic neurons or mGluR activation from the direct ascending projections provide the sustained spiking response throughout the minus phase. Reciprocal thalamocortical synapses back to Layer 6 might also assist in maintaining minus phase input drive. During this period, input from the sensory periphery is extremely downmodulated due to the relative quiescence of thalamic intrinsically bursting neurons. Downstream areas basically repeat this predictive process driven by their respective deep Layer 6 context, although they do also receive direct Layer 4 drive from the previous cortical area.

The plus phase is characterized by a shift of input from the local column to the sensory pe-




Figure 1.2: The LeabraTI computation

LeabraTI consists of a *minus* phase prediction about what is about to happen in the following *plus* phase sensory event. The minus phase is characterized by a sustained prediction from deep Layer 6 neurons based on lagged context from the previous moment. Input from the sensory periphery is extremely downmodulated or even blocked during the minus phase due to the relative quiescence of thalamic intrinsically bursting neurons. The plus phase is characterized by a rapid burst of information from the sensory periphery that drives also drives Layer 5 intrinsically bursting neurons to threshold so that they can integrate a snapshot of the current moment as context. Full feedforward and feedback-mediated processing is active in both phases so that predictions are not entirely dependent on deep Layer 6 neurons, but can also be driven by higher-level expectations.

riphery, driven by the destabilizing response from thalamic intrinsically bursting cells. The rapid influx of synchronous drive from the thalamus serves two purposes. First, it rapidly drives the propagation of responses through downstream areas by quickly driving neurons along the feedforward Layer 4 → Layer 2/3 → Layer 4 pathway (Felleman & Van Essen, 1991) to threshold, similar to the concept of “synfire chains” that simultaneously ignite entire pools of neurons to propagate waves of spikes back and forth through a network (Bruno & Sakmann, 2006; Wang, Spencer, Fellous, & Sejnowski, 2010; Tiesinga, Fellous, & Sejnowski, 2008). Second, it drives responses down

through the Layer 2/3 \rightarrow Layer 5 \rightarrow Layer 6 pathway, which also requires a large influx of thalamic drive (Beierlein, Fall, Rinzel, & Yuste, 2002). This rapid transmission between and within areas causes Layer 5b neurons across areas to burst with relatively little cross-area delay so that the context that the bursting captures represents information from a single moment in time.

Once thalamic bursting quiets, the input to the microcolumn shifts back to the newly integrated Layer 6 context. This prediction-sensation process repeats approximately every 100 ms corresponding to the brain's alpha rhythm, which is suggested to have strong thalamocortical and deep laminar sources (Lopes da Silva, 1991; Klimesch, Sauseng, & Hanslmayr, 2007; Palva & Palva, 2007; Lorincz et al., 2009; Hanslmayr et al., 2011; Buffalo, Fries, Landman, Buschman, & Desimone, 2011).

1.4 Relation to other frameworks

1.4.1 The Simple Recurrent Network (SRN)

LeabraTI is perhaps most closely related to the Simple Recurrent Network (SRN) (Elman, 1990; Servan-Schreiber et al., 1991), but with several key differences that are necessary due to the implementation using cortical circuits. The SRN is capable of learning temporal sequences by virtue of a “copy” operation (Figure 1.3) that represents an exact replica of the previous time step's intermediate units in separate group of context units. Thus, the SRN integrates two inputs at each time step – the previous time step's processed inputs and the current time step's unprocessed inputs. Predictions are made in a separate group of output units and prediction errors are backpropagated to learn the weights that allow the best prediction over items in the input sequence.

Like the SRN, LeabraTI uses discrete context units (deep Layer 6 neurons), but it does not implement a direct copy of sending units at each time step that is then used for predictive learning. LeabraTI is actually an inversion of the SRN formulation, where context is immediately sent through plastic integrative synapses and then through an additional input channel when computing the activation states of superficial (Layer 2 and 3) units (Figure 1.3). The LeabraTI formulation

is mathematically equivalent to the SRN, and provides a better fit with the circuitry of the microcolumn, which contains plastic synapses between Layer 5 and Layer 6 (primarily the 5a \rightarrow 5b integrative synapses). The inversion also allows a sustained, uncorrupted minus phase prediction that would not be possible if the directionality of the projections were reversed due to the massive numbers of afferent and efferent synapses on superficial neurons.




Figure 1.3: Relationship between the Simple Recurrent Network (SRN) and LeabraTI

The Simple Recurrent Network (SRN) projects a direct copy of Hidden units' state to context units that is then integrated at the next time step with the current input. LeabraTI, in contrast, projects first through Layer 5 \rightarrow Layer 6 plastic synapses triggering an immediate prediction that is then used additively in computing the input at the next time step. These operations are mathematically equivalent, but LeabraTI provides a better fit with the circuitry of the microcolumn.

One other difference between LeabraTI and the SRN is that the latter makes predictions using a separate group of output units. LeabraTI, in contrast, makes predictions using the same neural tissue that is used to represent sensory outcomes. This is a much more natural mechanism for predictive learning as it does not require a separate populations of neurons for predictions and sensations – it does, however, tradeoff the neural space savings with temporal interleaving of predictions and sensory outcomes, which could cause discretization artifacts and other temporal oddities in perception. There is increasing evidence of discretization at the ~ 10 Hz rate proposed by LeabraTI (VanRullen, Reddy, & Koch, 2005, 2006; Amano, Arnold, Takeda, & Johnston, 2008;

Busch et al., 2009; Sokoliuk & VanRullen, 2013; VanRullen & Koch, 2003; VanRullen & Dubois, 2011), which is discussed in the section following this one.

1.4.2 Predictive coding models

Another framework related to LeabraTI is the predictive coding model (Rao & Ballard, 1999; Friston, 2005; den Ouden, Kok, & de Lange, 2012). The fundamental claim of predictive coding models is that sensory events themselves are not explicitly encoded by neurons. Instead, prediction errors – the difference between predictions and sensory events – are what is encoded. The prediction error computation is generally repeated at each stage of processing with only the residual error being propagated downstream for further processing. Predictive coding models have gained traction due to this efficient “delta” code and due to their positing a distinct role for feedback projections in generating predictions. A recent description of predictive coding makes contact with the microcolumnar circuitry in the same way as LeabraTI. In their proposal, Bastos et al. (2012) suggested that deep infragranular neurons code prediction errors which ascend back to Layer 4 neurons or down to lower-level areas as feedback predictions.

The primary problem with predictive coding models is in their computation of prediction error. This computation requires feedback predictions to be inhibitory in nature so that they can be subtracted from the excitatory input, leaving only the residual error. This computation is not possible in cortical circuits since feedback via pyramidal axons is fundamentally excitatory. This issue could be resolved by inverting the excitatory signal, but feedback axons predominantly target other excitatory cells and the net postsynaptic potentials generated by their activation are also excitatory (Budd, 1998; Johnson & Burkhalter, 1996, 1997).

LeabraTI resolves this issue by computing predictions and sensory events explicitly in discrete phases, based on excitatory input with inhibitory competition to suppress spurious errors in coding (Wyatte, Herd, Mingus, & O’Reilly, 2012; Desimone & Duncan, 1995). Thus predictions are not fundamentally inhibitory, but excitatory due to being driven through Layer 4 neurons from deep context. This assumption is consistent with predictive remapping (Duhamel, Colby, & Gold-

berg, 1992; Nakamura & Colby, 2002; Cavanagh, Hunt, Afraz, & Rolfs, 2010) in cortex during which excitatory neurons fire in anticipation of a stimulus that is outside of their receptive field but will be entered upon making a saccade.

1.5 LeabraTI testable predictions

References

- Amano, K., Arnold, D. H., Takeda, T., & Johnston, A. (2008). Alpha band amplification during illusory jitter perception. *Journal of Vision*, *8*(10).
- Armstrong-James, M., Fox, K., & Das-Gupta, A. (1992). Flow of excitation within rat barrel cortex on striking a single vibrissa. *Journal of Neurophysiology*, *68*(4), 1345–1358.
- Bastos, A. M., Usrey, W. M., Adams, R. A., Mangun, G. R., Fries, P., & Friston, K. J. (2012). Canonical microcircuits for predictive coding. *Neuron*, *76*(4), 695–711.
- Beierlein, M., Fall, C. P., Rinzel, J., & Yuste, R. (2002). Thalamocortical bursts trigger recurrent activity in neocortical networks: Layer 4 as a frequency-dependent gate. *The Journal of Neuroscience*, *22*(22), 9885–9894.
- Belyusar, D., Snyder, A. C., Frey, H.-P., Harwood, M. R., Wallman, J., & Foxe, J. J. (2013). Oscillatory alpha-band suppression mechanisms during the rapid attentional shifts required to perform an anti-saccade task. *NeuroImage*, *65*, 395–407.
- Bruno, R. M., & Sakmann, B. (2006). Cortex is driven by weak but synchronously active thalamocortical synapses. *Science*, *312*(5780), 1622–1627.
- Budd, J. M. (1998). Extrastriate feedback to primary visual cortex in primates: A quantitative analysis of connectivity. *Proceedings of the Royal Society B*, *265*(1400), 1037–1044.
- Buffalo, E. A., Fries, P., Landman, R., Buschman, T. J., & Desimone, R. (2011). Laminar differences in gamma and alpha coherence in the ventral stream. *Proceedings of the National Academy of Sciences of the United States of America*, *108*(27), 11262–11267.
- Busch, N. A., Dubois, J., & VanRullen, R. (2009). The phase of ongoing EEG oscillations predicts visual perception. *The Journal of Neuroscience*, *29*(24), 7869–7876.
- Buxhoeveden, D. P., & Casanova, M. F. (2002). The minicolumn hypothesis in neuroscience. *Brain*, *125*(Pt 5), 935–951.
- Cavanagh, P., Hunt, A. R., Afraz, A., & Rolfs, M. (2010). Visual stability based on remapping of attention pointers. *Trends in Cognitive Sciences*, *14*(4), 147–153.
- Connors, B. W., Gutnick, M. J., & Prince, D. A. (1982). Electrophysiological properties of neocortical neurons in vitro. *Journal of Neurophysiology*, *48*(6), 1302–1320.
- da Costa, N. M., & Martin, K. A. C. (2010). Whose cortical column would that be? *Frontiers in Neuroanatomy*, *4*(16).

- den Ouden, H. E. M., Kok, P., & de Lange, F. P. (2012). How prediction errors shape perception, attention, and motivation. Frontiers in Psychology, 3(548).
- Desimone, R., & Duncan, J. (1995). Neural mechanisms of selective visual attention. Annual Review of Neuroscience, 18, 193–222.
- Douglas, R. J., & Martin, K. A. C. (2004). Neuronal circuits of the neocortex. Annual Review of Neuroscience, 27, 419–451.
- Duhamel, J. R., Colby, C. L., & Goldberg, M. E. (1992). The updating of the representation of visual space in parietal cortex by intended eye movements. Science, 255(5040), 90–92.
- Elman, J. L. (1990). Finding structure in time. Cognitive Science, 14(2), 179–211.
- Felleman, D. J., & Van Essen, D. C. (1991). Distributed hierarchical processing in the primate cerebral cortex. Cerebral Cortex, 1(1), 1–47.
- Foldiak, P. (1991). Learning invariance from transformation sequences. Neural Computation, 3(2), 194–200.
- Franceschetti, S., Guatteo, E., Panzica, F., Sancini, G., Wanke, E., & Avanzini, G. (1995). Ionic mechanisms underlying burst firing in pyramidal neurons: Intracellular study in rat sensorimotor cortex. Brain Research, 696(1–2), 127–139.
- Friston, K. (2005). A theory of cortical responses. Philosophical Transactions of the Royal Society B, 360(1456), 815–836.
- Gould, I. C., Rushworth, M. F., & Nobre, A. C. (2011). Indexing the graded allocation of visuospatial attention using anticipatory alpha oscillations. Journal of Neurophysiology, 105(3), 1318–1326.
- Hanslmayr, S., Gross, J., Klimesch, W., & Shapiro, K. L. (2011). The role of oscillations in temporal attention. Brain research reviews, 67(1–2), 331–343.
- Hirsch, J. A., & Martinez, L. M. (2006). Laminar processing in the visual cortical column. Current Opinion in Neurobiology, 16(4), 377–384.
- Horton, J. C., & Adams, D. L. (2005). The cortical column: A structure without a function. Philosophical Transactions of the Royal Society B, 360(1456), 837–862.
- Hubel, D. H., & Wiesel, T. N. (1977). Ferrier lecture. Functional architecture of macaque monkey visual cortex. Proceedings of the Royal Society B, 198(1130), 1–59.
- Hughes, S. W., Lorincz, M., Cope, D. W., Blethyn, K. L., Kekesi, K. A., Parri, H. R., Juhasz, G., & Crunelli, V. (2004). Synchronized oscillations at alpha and theta frequencies in the lateral geniculate nucleus. Neuron, 42(2), 253–268.
- Johnson, R. R., & Burkhalter, A. (1996). Microcircuitry of forward and feedback connections within rat visual cortex. The Journal of Comparative Neurology, 368(3), 383–398.
- Johnson, R. R., & Burkhalter, A. (1997). A polysynaptic feedback circuit in rat visual cortex. The Journal of Neuroscience, 17(18), 7129–7140.
- Jones, E. G. (2000). Microcolumns in the cerebral cortex. Proceedings of the National Academy of Sciences of the United States of America, 97(10), 5019–5021.

- Klimesch, W., Sauseng, P., & Hanslmayr, S. (2007). EEG alpha oscillations: The inhibition-timing hypothesis. Brain Research Reviews, 53(1), 63–88.
- Lee, C. C., & Sherman, S. M. (2009). Modulator property of the intrinsic cortical projection from layer 6 to layer 4. Frontiers in Systems Neuroscience, 3(3).
- Lopes da Silva, F. (1991). Neural mechanisms underlying brain waves: From neural membranes to networks. Electroencephalography and Clinical Neurophysiology, 79(2), 81–93.
- Lorincz, M. L., Crunelli, V., & Hughes, S. W. (2008). Cellular dynamics of cholinergically induced alpha (8–13 Hz) rhythms in sensory thalamic nuclei in vitro. The Journal of Neuroscience, 28(3), 660–671.
- Lorincz, M. L., Kekesi, K. A., Juhasz, G., Crunelli, V., & Hughes, S. W. (2009). Temporal framing of thalamic relay-mode firing by phasic inhibition during the alpha rhythm. Neuron, 63(5), 683–696.
- Lumer, E., Edelman, G., & Tononi, G. (1997). Neural dynamics in a model of the thalamocortical system. I. Layers, loops and the emergence of fast synchronous rhythms. Cerebral Cortex, 7(3), 207–227.
- Mathewson, K., Gratton, G., Fabiani, M., Beck, D., & Ro, T. (2009). To see or not to see: Prestimulus alpha phase predicts visual awareness. The Journal of Neuroscience, 29(9), 2725–2732.
- Mathewson, K. E., Prudhomme, C., Fabiani, M., Beck, D. M., Lleras, A., & Gratton, G. (2012). Making waves in the stream of consciousness: Entraining oscillations in EEG alpha and fluctuations in visual awareness with rhythmic visual stimulation. Journal of Cognitive Neuroscience, 24(12), 2321–2333.
- Mountcastle, V. B. (1997). The columnar organization of the neocortex. Brain, 120(Pt 4), 701–722.
- Nakamura, K., & Colby, C. L. (2002). Updating of the visual representation in monkey striate and extrastriate cortex during saccades. Proceedings of the National Academy of Sciences of the United States of America, 99(6), 4026–4031.
- O'Reilly, R. C., & Munakata, Y. (2000). Computational Explorations in Cognitive Neuroscience: Understanding the Mind by Simulating the Brain. Cambridge, MA: The MIT Press.
- O'Reilly, R. C., Munakata, Y., Frank, M. J., Hazy, T. E., & Contributors (2012). Computational Cognitive Neuroscience. Wiki Book, 1st Edition, URL: <http://ccnbook.colorado.edu>.
- Palva, S., & Palva, J. M. (2007). New vistas for alpha-frequency band oscillations. Trends in Neurosciences, 30(4), 150–158.
- Rao, R. P., & Ballard, D. H. (1999). Predictive coding in the visual cortex: A functional interpretation of some extra-classical receptive-field effects. Nature Neuroscience, 2(1), 79–87.
- Rockland, K. S., & Pandya, D. N. (1979). Laminar origins and terminations of cortical connections of the occipital lobe in the rhesus monkey. Brain Research, 179(1), 3–20.
- Rohenkohl, G., & Nobre, A. C. (2011). Alpha oscillations related to anticipatory attention follow temporal expectations. The Journal of Neuroscience, 31(40), 14076–14084.

- Schubert, D., Kotter, R., & Staiger, J. F. (2007). Mapping functional connectivity in barrel-related columns reveals layer- and cell type-specific microcircuits. Brain Structure & Function, 212(2), 107–119.
- Servan-Schreiber, D., Cleeremans, A., & McClelland, J. L. (1991). Graded state machines: The representation of temporal contingencies in simple recurrent networks. Machine Learning, 7(2–3), 161–193.
- Sherman, S., & Guillery, R. (2006). Exploring the Thalamus and Its Role in Cortical Function. Cambridge, MA: MIT Press.
- Silva, L. R., Amitai, Y., & Connors, B. W. (1991). Intrinsic oscillations of neocortex generated by layer 5 pyramidal neurons. Science, 251(4992), 432–435.
- Sokoliuk, R., & VanRullen, R. (2013). The flickering wheel illusion: When rhythms make a static wheel flicker. The Journal of Neuroscience, 33(33), 13498–13504.
- Thomson, A. M. (2010). Neocortical layer 6, a review. Frontiers in Neuroanatomy, 4(13).
- Thomson, A. M., & Lamy, C. (2007). Functional maps of neocortical local circuitry. Frontiers in Neuroscience, 1(1), 19–42.
- Tiesinga, P., Fellous, J.-M., & Sejnowski, T. J. (2008). Regulation of spike timing in visual cortical circuits. Nature Reviews Neuroscience, 9(2), 97–107.
- VanRullen, R., & Dubois, J. (2011). The psychophysics of brain rhythms. Frontiers in Psychology, 2(203).
- VanRullen, R., & Koch, C. (2003). Is perception discrete or continuous? Trends in Cognitive Sciences, 7(5), 207–213.
- VanRullen, R., Reddy, L., & Koch, C. (2005). Attention-driven discrete sampling of motion perception. Proceedings of the National Academy of Sciences of the United States of America, 102(14), 5291–5296.
- VanRullen, R., Reddy, L., & Koch, C. (2006). The continuous wagon wheel illusion is associated with changes in electroencephalogram power at approximately 13 Hz. The Journal of Neuroscience, 26(2), 502–507.
- Wang, H.-P., Spencer, D., Fellous, J.-M., & Sejnowski, T. J. (2010). Synchrony of thalamocortical inputs maximizes cortical reliability. Science, 328(5974), 106–109.
- Wyatte, D., Herd, S., Mingus, B., & O'Reilly, R. (2012). The role of competitive inhibition and top-down feedback in binding during object recognition. Frontiers in Psychology, 3, 182.